

EpiKG: End-to-End Epistemic Preservation in Clinical Knowledge Graphs for Assertion-Aware Retrieval-Augmented Generation

Anonymous Author(s)

Abstract

Clinical NLP extracts assertion status (negated, uncertain, family-history, etc.), but these labels are often dropped before retrieval, causing downstream QA to treat epistemically distinct facts as equivalent. We call this failure mode the *epistemic propagation gap*. We introduce EpiKG, a clinical KG-RAG pipeline that preserves a 7-value assertion schema and tri-temporal metadata on graph edges, then uses intent-aware routing to select question-type-specific traversals.

On ClinicalBench (400 questions, 43 MIMIC-IV patients; Claude Opus 4.6), intent-aware KG-RAG (C4g) reaches 69.0% versus 21.8% for LLM-alone (+47.2 pp). Ablation isolates the mechanism: structured retrieval contributes +28.2 pp (C1→C3), adding assertions without routing hurts (−3.2 pp, C3→C4), and intent-aware routing recovers and amplifies gains (+22.2 pp, C4→C4g). Thus, assertion metadata is necessary but not sufficient. Long-context concatenation (C6) reaches 39.0% but fails on historical and sequence (0% each), showing that context scale alone does not solve longitudinal reasoning.

Cross-model validation (Opus 4.6, GPT-OSS 20B, MedGemma 27B) shows consistent C1→C4g gains (+31.5 to +47.2 pp), with model-dependent magnitude. In a blinded physician pilot ($n = 30$), C4g answers were rated more correct (81% vs. 29%), safer (88% vs. 50%), and more helpful (81% vs. 29%). Physician review also found 40% errors in auto-generated gold labels, motivating claim calibration and multi-evaluator reporting. Overall, we show that category×condition interactions, rather than aggregate accuracy alone, are the critical lens for evaluating clinical KG-RAG.

1 Introduction

A physician documents “patient denies chest pain, sister had MI at 45, will consider statin if lipids remain elevated”—a single sentence encoding four epistemic states: a negated symptom, a family-history event, a hypothetical action, and an implicit present condition. Clinical NLP can detect these assertions accurately [1, 2], yet when a RAG system ingests the same note, assertion context is flattened: negation is lost, family attribution collapses onto the patient, and downstream reasoning may conclude the patient *has* chest pain and *had* an MI. We call this the *epistemic propagation gap*—the systematic loss of assertion status as clinical information moves from extraction through storage to retrieval.

The gap persists because, to our knowledge, prior KG-RAG systems do not propagate assertion labels beyond the extraction stage into graph-augmented retrieval. Clinical NLP pipelines (cTAKES, John Snow Labs) detect assertions, but do not carry them as edge attributes into knowledge graph construction or retrieval. OMOP (Observational Medical Outcomes Partnership) excludes negated conditions from `CONDITION_OCCURRENCE` [3]; FHIR (Fast Healthcare Interoperability Resources) provides `verificationStatus` for `Condition` resources only, covering a fraction of the clinical data model. Graph-augmented RAG systems—GraphRAG [4], GFM-RAG [5], KARE [6], Medical-Graph-RAG [7]—build KGs that discard the metadata distinguishing “patient has diabetes” from “rule out diabetes.” Two KG surveys [8, 9] do not identify assertion preservation as a concern. A parallel *temporal integration gap* exists: clinical events span three time dimensions (valid time, transaction time, NLP-asserted temporality), yet existing systems model at most two [10, 11]. Allen’s

interval algebra [12] has been formalized as annotation ontologies [13] and modal operators [14], but never as first-class KG edge attributes that participate in retrieval.

We present EpiKG, a clinical KG system that closes both gaps by preserving epistemic status end-to-end from NLP extraction through OMOP mapping, fact construction, KG materialization, and graph-augmented retrieval. Our contributions:

1. **Epistemic preservation invariant.** We formalize the requirement that assertion status remain recoverable across extraction, mapping, storage, and retrieval; derive an information-theoretic loss bound for assertion-blind pipelines (Section 3.5, Appendix B); and implement a seven-value assertion schema carried end-to-end.
2. **Tri-temporal epistemic KG.** We represent valid time, transaction time, and NLP-asserted temporality jointly on edges, with Allen-style temporal relations as first-class retrieval features, enabling longitudinal reasoning over multi-admission records.
3. **Ablation design that isolates mechanism.** We introduce C4 (assertions without intent routing) between C3 and C4g, and include C6/C7 as bookend baselines. This separates the effect of structured retrieval, assertion metadata, and routing policy instead of reporting a single bundled gain.
4. **Multi-evaluator and cross-model validation.** We evaluate on two benchmarks (ClinicalBench, SliceBench), three models, and three evaluator views (deterministic keyword, LLM-as-judge, blinded physician pilot), reporting where estimates agree, where they diverge, and which conclusions are robust to evaluator choice.
5. **A falsifiable empirical claim.** The key claim is not that “KG-RAG is always better,” but that epistemic metadata improves QA *only when retrieval uses it*. Empirically, assertions without routing can hurt ($C3 \rightarrow C4$), while routing unlocks gains ($C4 \rightarrow C4g$), especially on hard longitudinal categories.

Our central research question is: *under what conditions does structured epistemic context help, hurt, or break even relative to unstructured retrieval and no retrieval?* The answer is interactional rather than uniform: gains depend on question type, routing policy, and model. Across experiments, intent-aware KG-RAG improves all nine ClinicalBench categories, with largest effects on cross-admission longitudinal tasks; assertions without routing can degrade temporal categories; and brute-force long context does not recover these failures despite a large token window. These results motivate a narrower and more testable takeaway: category $\times$ condition interactions are the appropriate evaluation target for clinical KG-RAG, while aggregate accuracy alone obscures mechanism (Section 5).

2 Related Work

We organize prior work along four axes and synthesize the gaps that EpiKG addresses (Table 9, Appendix).

Medical RAG systems. Graph-augmented retrieval has emerged as a leading paradigm for medical QA. GraphRAG [4] introduced community-based summarization; GFM-RAG [5] trained an 8M-parameter graph foundation model across 60 KGs (14M+ triples); KARE [6] adapted community retrieval to clinical decision support; DoctorRAG [15] emulates physician reasoning; MedRAG [16] constructs hierarchical diagnostic KGs; Medical-Graph-RAG [7] links documents via triple graphs. Concurrent work addresses EHR-specific retrieval: EHR-RAG [17] evaluates retrieval strategies over structured EHR tables, and MediGRAF [18] builds medical KGs via multi-agent extraction. Yet no published medical RAG system models assertion status or temporal context; GFM-RAG and KARE build population-level graphs rather than patient-level graphs from clinical text, and a recent survey [19] confirms no system incorporates epistemic metadata into the retrieval graph.

Clinical KG construction. Multi-LLM KG-RAG [20] uses multi-agent LLM prompting with schema-constrained extraction for oncology, modeling uncertainty via entropy but not tracking assertion classes on edges. AutoRD [21] applies GPT-4 to rare disease literature; RECAP-KG [22]




Figure 1: The EpiKG pipeline. The assertion label α (highlighted) is preserved as a first-class attribute from NLP extraction through OMOP mapping, fact construction, KG materialization, and assertion-aware GraphRAG.

handles messy GP notes; FHIR-Ontop-OMOP [23] generates virtual KGs from OMOP data. All share a common limitation: assertion status is not propagated into the final graph.

Assertion detection. NegEx [24] introduced trigger-based negation detection; ConText [25] extended it with temporality and experienter. The i2b2/VA shared task [1] formalized a six-class taxonomy; Gul et al. [2] fine-tuned LLMs to 0.962 accuracy. All treat assertion detection as a terminal annotation task: labels are not carried into knowledge graphs or retrieval systems.

Assertion detection is necessary but not sufficient: labels must also be *temporally situated*—the same condition may be present in one encounter and absent in the next.

Temporal knowledge graphs. MedTKG [10] constructs temporal KGs with time-stamped snapshots (event time only); Graphiti [11] implements bitemporal edges [26] but lacks clinical ontology alignment; TEO [13] and TCL [14] formalize Allen’s relations [12] as annotation ontologies or modal operators rather than KG edge attributes. Uncertainty quantification in biomedical KGs [27] has not been connected to assertion taxonomies. Two structural gaps emerge: an *epistemic propagation gap* (assertion-detecting systems do not persist labels into KGs) and a *temporal integration gap* (temporal formalisms operate as annotation layers, not as KG edge attributes participating in retrieval). EpiKG closes both by carrying assertion and temporal metadata as first-class properties through every pipeline stage.

3 System Design

3.1 Overview

EpiKG transforms unstructured clinical narratives into an assertion-aware, temporally grounded knowledge graph exposed through graph-augmented retrieval. The pipeline comprises six stages: (1) document ingestion, (2) NLP extraction, (3) OMOP concept normalization [3], (4) clinical fact construction with epistemic metadata, (5) KG materialization with shared concept nodes, and (6) assertion-aware retrieval. The key invariant is that *epistemic assertion status* is a first-class attribute at every stage, not an NLP annotation discarded after extraction (Figure 1). Clinical concepts are stored once as shared global nodes; patient-specific semantics (assertion status, temporality, confidence) reside on edges, enabling cohort queries via edge scans.

3.2 Epistemic Assertion Schema

Clinical notes abound with qualified statements—“*no evidence of pneumonia*,” “*possible CHF*,” “*mother had breast cancer*”—yet standard representations discard these qualifiers. OMOP excludes negated conditions entirely [3]; FHIR limits assertion metadata to conditions. EpiKG defines a seven-value assertion taxonomy:

$$\alpha \in \{\text{Pres., Abs., Poss., Cond., Hypo., Fam.Hx., Hist.}\} \quad (1)$$

This extends the i2b2 (Informatics for Integrating Biology and the Bedside) six-class taxonomy [1] by separating HISTORICAL from FAMILY_HISTORY, a clinically important distinction that prior work conflates (definitions in Appendix Q). A rule-based classifier (122 scope-aware trigger patterns) assigns α with an associated confidence score; we chose rule-based classification for deterministic behavior and interpretability, though fine-tuned LLM classifiers [2] could be substituted. Intrinsic classifier evaluation is deferred; the end-to-end ClinicalBench ablation (C3→C4g) serves as a functional test of assertion propagation. The label is carried through every stage: from ExtractedMention through ensemble merging, ClinicalFact records, KG edge properties, and retrieval scoring. To our knowledge, no prior system [2, 28] maintains this end-to-end invariant.

3.3 Tri-Temporal Knowledge Graph Model

Clinical events unfold across multiple time dimensions that prior systems model incompletely [10, 11]. Bi-temporal databases [26] distinguish valid time from transaction time; EpiKG adds a third dimension (NLP-asserted temporality) and stores all three on every KG edge:

Definition 1 (Tri-Temporal Edge). *An edge $e = (s, p, o, \alpha, \tau_v, \tau_t, \tau_a, r, c)$ where:*

- s, o are source and target nodes; p is the predicate (one of 24 edge types);
- $\alpha \in \mathcal{A}$ is the assertion label (Eq. 1);
- $\tau_v = (\text{event_date}, \text{valid_from}, \text{valid_to})$ is the **valid time** interval—when the relationship held in the real world;
- $\tau_t = (\text{recorded_at}, \text{doc_date}, \text{created_at})$ is the **transaction time**—when it was recorded;
- $\tau_a \in \{\text{CURRENT}, \text{PAST}, \text{FUTURE}\}$ is the **NLP-asserted temporality**;
- $r \in \mathcal{R}$ is an Allen interval algebra relation [12]; $c \in [0, 1]$ is temporal confidence.

$\mathcal{R}$ comprises seven of Allen’s 13 canonical relations [12] (merging symmetric pairs like Before/Meets) plus CONCURRENT and UNKNOWN (nine total; full mapping in Appendix R). Unlike TEO [13] and TCL [14], which use Allen’s relations as annotation-ontology classes or modal operators, EpiKG stores r and c directly on materialized edges, enabling efficient temporal filtering during traversal.

3.4 Assertion-Aware Graph-Augmented Retrieval

Given a clinical question q and patient π , the retrieval pipeline: (1) extracts concepts from q via NLP + OMOP enrichment; (2) traverses 2–3 hops via bidirectional breadth-first search (BFS) over patient KG edges and OMOP vocabulary relationships (20M+ edges) via PostgreSQL common table expressions (CTEs) (Appendix K), pruning edges with $c < 0.3$; (3) groups edges into four temporal views (event timeline, current state, historical, conflicts); (4) retrieves matching guideline sections; (5) scores edges by:

$$\text{score}(e) = c(e) + \mathbb{I}[\text{type}(e) \in Q_{\text{rel}}] \cdot 0.2 + \mathbb{I}[\tau_a(e) = \text{CURRENT}] \cdot 0.1 \quad (2)$$

where the bonuses for question-relevant edge types (0.2) and current temporality (0.1) were set by manual tuning on a development set of 20 questions. The surviving subgraph is serialized into structured text preserving assertion labels (e.g., “ABSENT: pneumonia”); and (6) composes graph evidence, temporal context, guidelines, and source documents into a single prompt.

3.5 Formal Epistemic Preservation

We formalize the epistemic invariant maintained by the EpiKG pipeline, building on the ontological principle that aboutness must be preserved across transformations [29]. The *epistemic state* of a clinical mention m is the tuple $e(m) = (c, \alpha, \xi, \tau)$: OMOP concept, 7-value assertion

Table 1: ClinicalBench conditions. C1–C4g form the ablation ladder; C6 and C7 are bookend baselines.

ID	Condition	Retrieval	Assertion	Temporal
C1	LLM Alone	None	None	None
C2	+ Vanilla RAG (TF-IDF)	TF-IDF doc chunks	None	None
C2b	+ Vanilla RAG (dense)	Contriever doc chunks	None	None
C3	+ KG-RAG (no assertions)	Graph + Doc	None	None
C4	+ Epistemic KG-RAG	Graph + Doc	Full (7-class)	Tri-temporal
C4g	+ Intent-Aware KG-RAG	Graph + Doc (type-specific)	Full (7-class)	Tri-temporal
C6	Long Context	All notes	None	None
C7	Deterministic KG	KG lookup (no LLM)	Full	Tri-temporal

label, experiencer ($\{\text{PATIENT, FAMILY}\}$), and temporality ($\{\text{CURRENT, PAST, FUTURE}\}$); see Table 8 (Appendix) for notation summary. A pipeline *epistemically preserves* m if the output epistemic state equals the extraction state at every stage. We formalize (Appendix B) that an assertion-blind pipeline collapses the assertion distribution to a degenerate point mass, incurring entropy loss $\Delta H = H(A_c^\pi) \geq 0$ [30], and that assertion-faithful accuracy is bounded by $1 - f_{np}(c)$, where f_{np} is the fraction of non-present assertions. The empirical consequences are measured via category-stratified accuracy in Section 5.

3.6 Intent-Aware Retrieval (C4g)

The multi-hop BFS described above treats all questions uniformly, yet different clinical question types require fundamentally different graph operations. A *change* question requires cross-admission set differencing; a *current-state* question needs filtering to the most recent valid edges; a *historical* question must recover resolved conditions.

Intent classifier. A lightweight rule-based classifier maps each question to one of four intents—CHANGE, CURRENT_STATE, HISTORICAL, or DEFAULT—using question metadata with keyword-pattern fallback (Algorithm 1, Appendix S).

Change retrieval. For $\iota = \text{CHANGE}$, edges are partitioned by admission (`hadm_id`). The module computes set differences across admission pairs (k, k') —concepts *added* ($C_k \setminus C_{k'}$), *removed* ($C_k \setminus C_{k'}$), and *continued* ($C_k \cap C_{k'}$)—and returns structured evidence.

Current-state and historical retrieval. For CURRENT_STATE, the module filters to edges with $\tau_a = \text{CURRENT}$ or open-ended validity, deduplicates by concept, and emits “NOT FOUND” for missing concepts. For HISTORICAL, it selects edges with $\tau_a = \text{PAST}$ and augments with admission-based inference (concepts in earlier admissions but absent from the latest are labeled “resolved”). Each intent triggers a type-specific prompt template that structures evidence according to question semantics (Appendix U).

4 Benchmark Design

Existing clinical QA benchmarks evaluate factual recall (MedQA [31]) or agent task completion [32] but do not test handling of epistemic qualifiers or reasoning over longitudinal records of varying complexity. We introduce two complementary benchmarks built on MIMIC-IV [33].

4.1 ClinicalBench: Assertion-Sensitive Clinical QA

ClinicalBench comprises 400 gold-standard questions over 43 MIMIC-IV patients in two tasks: **Task A** (200 questions, negation-aware retrieval) and **Task B** (200 questions, temporal reasoning). Questions span 9 categories: *negation*, *conditional*, *uncertainty*, *family_history*, *sequence*, *current_state*, *duration*, *historical*, and *change*. Gold-standard answers were created by a single physician (board-certified emergency medicine) through manual chart review against source MIMIC-IV documents. No inter-annotator agreement was computed for gold-standard creation; the physician adjudication pilot (Section 4.4) addresses this limitation. Four ablation conditions progressively add system components (Table 1). Two bookend baselines—C6 (all notes concatenated) and C7 (deterministic KG lookup, no LLM)—bracket the design space.

Table 2: ClinicalBench cohort demographics ($n = 43$ MIMIC-IV patients).

Characteristic	Category	n (%)
Sex	Female	26 (60.5%)
	Male	17 (39.5%)
Age (anchor)	18–29	2 (4.7%)
	30–49	13 (30.2%)
	50–64	13 (30.2%)
	65–79	9 (20.9%)
	80+	6 (14.0%)
Race/Ethnicity	White	33 (76.7%)
	Black/African American	6 (14.0%)
	Asian	1 (2.3%)
	Other/Unknown	3 (7.0%)
Insurance	Medicare	19 (44.2%)
	Private	15 (34.9%)
	Medicaid/Other	9 (20.9%)

Cohort demographics. The 43 ClinicalBench patients (Table 2) are drawn from a single academic medical center (Beth Israel Deaconess, Boston). The cohort skews female (60.5%), White (76.7%), and Medicare-insured (44.2%), reflecting the source institution’s patient population. Subgroup accuracy analyses are underpowered given the small cohort; we report them transparently in Appendix X rather than drawing subgroup conclusions.

Endpoints. The **pre-registered primary endpoint** is the hard longitudinal subset (*change + current_state + historical*, $n = 130$): C4g vs. C1 aggregate accuracy. The **secondary endpoint** is full 400-question ClinicalBench, C4g vs. C1. **Diagnostic endpoints** are per-category C1→C4g accuracy deltas and the C3→C4→C4g decomposition isolating assertion metadata (C3→C4) from intent routing (C4→C4g).

What each transition isolates. C1→C2/C2b isolates the *document retrieval effect*: C2 adds TF-IDF and C2b adds Contriever dense retrieval over clinical notes, analogous to existing medical RAG systems (DoctorRAG, MedRAG). C2/C2b→C3 isolates the *graph structure effect*: both retrieve documents, but C3 replaces document-level retrieval with graph traversal—testing whether structured retrieval outperforms vanilla retrieval. C3→C4 isolates the *assertion metadata effect*: C4 adds 7-class assertions and tri-temporal metadata but uses uniform BFS (no intent routing). C4→C4g isolates the *intent routing effect*: C4g adds question-type-specific graph traversal and assertion-aware filtering on top of C4’s metadata. C4g vs. C6 isolates the *structured-vs-unstructured retrieval effect* on the same questions.

4.2 SliceBench: Complexity-Stratified Longitudinal QA

SliceBench tests the hypothesis that KG-augmented retrieval provides increasing benefit as patient record complexity grows. We select 6 MIMIC-IV patients in three tiers: **Tier A** (2 patients, 1–2 encounters), **Tier B** (2 patients, 5–10), and **Tier C** (2 patients, 15+). Each patient receives 24 questions (including hard longitudinal categories: cross-encounter medication timelines, problem list reconciliation, causal chain tracing), yielding 144 total. Five conditions (B0–B4) form a monotone context progression from parametric-only to full system (Appendix J). The critical comparison is B2→B3: both have the same documents, but B3 adds structured KG context with assertion metadata and temporal relations.

4.3 Evaluation Protocol

ClinicalBench uses a *deterministic keyword evaluator* (v2) that performs exact word-boundary matching of gold-standard assertion and temporal keywords against model answers (keyword lists in Appendix P). The v2 evaluator includes an *abstention detection gate*: when a model responds that information is unavailable in the notes (e.g., “cannot determine,” “not mentioned”) rather than making a clinical claim, the answer is scored as incorrect regardless of keyword matches. This prevents false positives where refusal language incidentally matches gold-standard patterns. Sequence and change categories additionally require domain-specific keywords (ordering terms and change

Table 3: ClinicalBench primary results (400 questions, Claude Opus 4.6, keyword evaluator v2 with abstention detection). Ablation ladder (C1–C4g) progressively adds components; C4 isolates assertion metadata from intent routing; C6 and C7 are bookend baselines. Sig: * denotes BCa CI excluding zero. Best in **bold**.

	Condition	Accuracy	Δ vs C1
C1	LLM Alone	21.8%	—
C2	+ Vanilla RAG (TF-IDF)	52.2%	+30.5 pp *
C2b	+ Vanilla RAG (dense)	50.7%	+29.0 pp *
C3	+ KG-RAG (no assertions)	50.0%	+28.2 pp *
C4	+ Assertions (no routing)	46.8%	+25.0 pp *
C4g	+ Intent-Aware KG-RAG	69.0%	+47.2 pp *
C6	Long Context (all notes)	39.0%	+17.2 pp
C7	Deterministic KG (no LLM)	3.8%	−18.0 pp

terms, respectively). This evaluator is fully reproducible (no stochastic LLM judge) and is used for all ClinicalBench results. The primary answering model is Claude Opus 4.6; cross-model validation uses MedGemma 27B and GPT-OSS 20B under the same evaluator.

SliceBench uses an *LLM-as-judge* protocol with separated answering and judging models to prevent self-evaluation bias [34]: Claude Sonnet 4.5 generates answers and Claude Opus 4.6 judges correctness.

Statistical reporting. We report BCa bootstrap 95% CIs ($n = 2,000$, seed 42) [35] on both benchmarks. Question-level resampling is the primary analysis; patient-level cluster bootstrap (resampling the 43 patients with replacement, including all their questions) is reported as a sensitivity analysis to account for within-patient correlation (Appendix Table 17). McNemar’s test [36] provides a complementary non-parametric paired comparison; Benjamini–Hochberg FDR correction is applied to the family of McNemar tests (all remain significant after correction; the non-significant C3→C4 comparison, $p_{BH} = 0.37$, is unaffected). ClinicalBench CIs appear on primary deltas (overall C4g–C1 and hard longitudinal subset); SliceBench CIs appear on all conditions (Table 6). A safety score with asymmetric weighting ($w = 2.0$ for false-positive assertion errors) captures clinical risk (Appendix D).

4.4 Physician Validation Protocol

To validate automated scoring, a board-certified physician (A.S.) reviewed a blinded, stratified pilot sample ($n = 30$ ClinicalBench items across all 9 categories and both conditions, C1 vs. C4g). The physician-rated C4g–C1 delta (+52 pp) exceeded the auto-evaluator delta (+47 pp), indicating the headline effect is conservative (Section 5.5). A second independent reviewer is planned; protocol details are in Appendix O.

5 Results

We evaluate EpiKG through two benchmarks: ClinicalBench tests assertion-sensitive and longitudinal reasoning, while SliceBench measures longitudinal QA across patient complexity tiers. ClinicalBench uses a deterministic keyword evaluator for reproducibility; SliceBench uses LLM-as-judge evaluation (Section 4.3). Claude Opus 4.6 is the primary answering model; MedGemma 27B and GPT-OSS 20B¹ provide cross-model validation. Both benchmarks report bias-corrected and accelerated (BCa) bootstrap 95% CIs ($n = 2,000$, seed 42) [35], resampled at the question level.

5.1 ClinicalBench: Primary Ablation

Intent-aware KG-RAG (C4g) reaches 69.0% (+47.2 pp vs. C1; 95% BCa CI: [+40.9 pp, +52.8 pp]; patient-level cluster bootstrap: [+40.2 pp, +52.5 pp]; McNemar $\chi^2 = 157.1$, $p < 10^{-6}$). Without retrieval, the model mostly abstains or produces vague responses that fail keyword matching; with structured retrieval, it produces substantive correct answers. The bookend baselines bracket the design space. C6 (all notes concatenated) scores 39.0%, outperforming C1 (+17.2 pp) because it

¹GPT-OSS 20B is an open-source general-purpose model run locally via Ollama (4-bit quantized). We use the label GPT-OSS to avoid implying endorsement; the reproducibility package includes model identifiers and raw predictions.

provides clinical text. However, C6 fails on historical (0%) and sequence (0%), consistent with the “lost in the middle” phenomenon when context exceeds what the model can attend to. C7 (deterministic KG, no LLM) achieves 3.8%—only negation detection works (13.6%) via stored assertion metadata—confirming that an LLM is essential for reasoning over graph evidence.

Ablation decomposition: retrieval vs. assertions vs. intent routing. C2 (TF-IDF document retrieval) and C2b (dense retrieval via Contriever) score comparably—52.2% and 50.7%, respectively ($\Delta = -1.5$ pp; 95% BCa CI: [-6.6 pp, +3.5 pp]; McNemar $p = 0.62$)—confirming that the retrieval method (keyword vs. dense) does not explain the C2→C4g gap. C3 (graph-augmented retrieval without assertions) also scores similarly (50.0%), demonstrating that graph structure alone does not improve aggregate accuracy over either retrieval baseline. However, per-category patterns diverge: C3 excels on current state (54% vs. 32%) where graph structure helps identify active conditions, while C2 outperforms C3 on change (43% vs. 17%) where raw text explicitly mentions medication changes.

The new C4 condition isolates assertion metadata from intent routing: C4 adds epistemic annotations (negation, uncertainty, temporality) to graph edges but uses uniform BFS traversal without question-type-specific routing. C4 scores 46.8%—*below* C3 (50.0%)—with the C3→C4 delta of -3.2 pp not reaching significance (95% BCa CI: [-10.0 pp, +3.2 pp]; McNemar $p = 0.37$). This reveals that assertion metadata without intent routing actually degrades performance: the additional edge properties increase evidence verbosity and confuse uniform BFS traversal. The entire C3→C4g gain (+19.0 pp; CI: [+13.8 pp, +24.0 pp]) is attributable to intent-aware routing (C4→C4g: +22.2 pp; CI: [+16.0 pp, +28.0 pp]; McNemar $\chi^2 = 41.4$, $p < 10^{-10}$).

Per-category decomposition clarifies the mechanism: assertions without routing help on assertion-sensitive categories (negation: +24.5 pp, uncertainty: +15.0 pp) but *hurt* on temporal categories (current state: -24.0 pp, historical: -30.0 pp, sequence: -47.5 pp). Intent routing reverses these losses and amplifies the gains—most dramatically on change (+90.0 pp), sequence (+55.0 pp), and historical (+48.0 pp). The finding that assertions are necessary but not sufficient—and that their value is unlocked specifically by routing—validates the two-component architecture.

C2/C2b serve as the closest analogs to existing medical RAG systems (DoctorRAG, MedRAG), which use document retrieval without assertion-aware graph structure; the +18.2 pp C2b→C4g gap (CI: [+12.3 pp, +23.5 pp]; McNemar $p < 10^{-9}$) quantifies the value of epistemic KG-RAG over modern dense retrieval.

Evaluator hierarchy. We report three evaluators to bound the true effect size. The *primary* evaluator is the deterministic keyword matcher (v2 with abstention detection), chosen for reproducibility: it is fully deterministic and produces the same score on every run. Two *sensitivity analyses* probe its biases. First, a blinded physician pilot (Section 5.5) found the keyword evaluator too strict in 43% of cases (vs. too lenient in 10%); the physician-rated C4g–C1 delta (+52 pp) exceeds the keyword delta (+47 pp), indicating the headline effect is conservative. Second, an LLM-as-judge evaluation (Appendix V) achieves higher physician concordance (56.7% vs. 46.7%) and yields a C4g–C1 delta of +28.5 pp—lower because the judge credits partial answers and appropriate C1 abstentions that the keyword evaluator penalizes. The three deltas (+28.5 pp, +47.2 pp, +52 pp) bound the effect: the keyword evaluator understates C1 (conservative bias), while the physician-rated delta may overstate it given small n . All three agree that C4g substantially outperforms C1; the disagreement concerns magnitude, not direction.

5.2 Category × Condition Interaction

Table 4 and Figure 2 reveal a strong *category × condition interaction*: intent-aware KG-RAG improves all 9 categories over C1, with every delta exceeding +35 pp. Per-category results are diagnostic and exploratory; no multiplicity correction is applied. The preregistered primary endpoint is the hard longitudinal subset (change + current state + historical). Under the v2 evaluator with abstention detection, C1 drops sharply because the model without retrieval context frequently abstains (“insufficient information in the notes”) rather than attempting an answer.

Where intent-aware routing excels. The largest gains occur on categories requiring cross-admission synthesis: change (+93 pp), sequence (+55 pp), historical (+54 pp), family history

Table 4: ClinicalBench per-category accuracy (%) across conditions (Claude Opus 4.6, evaluator v2). C2b replaces TF-IDF with Contriever dense retrieval; C4 adds assertion metadata to C3 without intent routing; C4g adds intent-aware routing. Best per row in **bold**.

Category	n	C1	C2	C2b	C3	C4	C4g	C4g-C1
Negation	110	45.5	70.9	75.5	65.5	90.0	80.9	↑ 35.5 pp
Conditional	20	0.0	35.0	50.0	30.0	50.0	45.0	↑ 45.0 pp
Uncertainty	40	12.5	35.0	32.5	37.5	52.5	50.0	↑ 37.5 pp
Family hist.	30	3.3	43.3	43.3	43.3	50.0	56.7	↑ 53.3 pp
Sequence	40	10.0	60.0	47.5	57.5	10.0	65.0	↑ 55.0 pp
Current st.	50	26.0	32.0	32.0	54.0	30.0	70.0	↑ 44.0 pp
Duration	30	30.0	66.7	66.7	60.0	46.7	66.7	↑ 36.7 pp
Historical	50	6.0	48.0	46.0	42.0	12.0	60.0	↑ 54.0 pp
Change	30	6.7	43.3	20.0	16.7	10.0	100.0	↑ 93.3 pp
Overall	400	21.8	52.2	50.7	50.0	46.8	69.0	↑ 47.2 pp




Figure 2: ClinicalBench per-category accuracy across six conditions (C6, C1, C2, C3, C4, C4g; Claude Opus 4.6, evaluator v2). C6 (pink dash-dot) fails on historical and sequence (both 0%); C1 (gray dotted) falls below 50% on 7 of 9 categories; C2 (light blue) and C3 (orange dashed) score similarly overall (~50–52%) but diverge per category; C4 (firebrick, short-dash) adds assertions without routing and scores 46.8%—below C3—demonstrating that assertions alone are insufficient; C4g (dark solid) adds intent routing, unlocking the value of epistemic annotations.

(+53 pp), conditional (+45 pp), and current state (+44 pp). On the hard longitudinal subset (change + current state + historical, $n = 130$), C4g outperforms C1 by +59.2 pp (95% BCa CI: [+49.2 pp, +67.7 pp]; patient-level: [+48.6 pp, +69.4 pp]). All nine category improvements exceed +35 pp (Appendix Figure 4, right panel); the left and middle panels decompose the C3→C4g gain into assertion

Table 5: ClinicalBench cross-model results (400 questions, keyword evaluator v2). All three models benefit from KG-RAG; deltas vary by model.

	Model	C1	C4g	Δ
Commercial	Claude Opus 4.6	21.8%	69.0%	+47.2 pp
Open	GPT-OSS 20B	20.5%	58.0%	+37.5 pp
Medical	MedGemma 27B	26.2%	57.8%	+31.5 pp

Table 6: SliceBench overall results (144 questions, Claude Sonnet 4.5, judged by Opus 4.6). Sig: * denotes CI excluding zero.

	Condition	Score	95% CI	Δ vs B0
B0	LLM Alone	49.9%	[43.6%, 56.1%]	—
B1	Latest Note	73.4%	[67.8%, 78.7%]	+23.5 pp
B2	All Notes RAG	84.7%	[80.2%, 88.8%]	+34.8 pp *
B3	KG-RAG	86.9%	[82.8%, 90.8%]	+37.0 pp *
B4	Full System	87.6%	[83.7%, 91.2%]	+37.8 pp *

321 metadata (C3→C4) and intent routing (C4→C4g), showing that assertions alone help assertion-
 322 sensitive categories but degrade temporal ones. C6 fails catastrophically on historical (0%) and
 323 sequence (0%) despite Opus’s 200K window, confirming that long context cannot substitute for
 324 structured retrieval on these categories.

325 5.3 Cross-Model Validation

326 All three models benefit substantially from KG-RAG (Table 5), with a consistent trend: models
 327 with lower C1 baselines gain more from structured retrieval. Opus and GPT-OSS start with similar
 328 C1 baselines (21.8% and 20.5%) but Opus gains the most (+47.2 pp), while MedGemma—the only
 329 model fine-tuned on medical corpora—starts highest (26.2%) yet gains least (+31.5 pp). The pat-
 330 tern is consistent with a knowledge-compensator effect, though with only three models the trend is
 331 suggestive rather than conclusive.

332 All three models improve on all 9 categories under C4g (per-category cross-model breakdowns
 333 in Appendix Table 10). Intent-aware KG-RAG thus functions as a *medical knowledge compensator*:
 334 the structured epistemic context partially substitutes for parametric medical knowledge, providing
 335 the largest benefit to general-purpose models that need it most.

336 5.4 SliceBench

337 The incremental KG layer (B2→B3) contributes +2.2 pp overall (CI: [-1.5 pp, +5.9 pp]), not
 338 reaching aggregate significance, but tier-stratified estimates show a complexity-dependent trend:
 339 +0.6 pp (Tier A), +1.0 pp (Tier B), +5.0 pp (Tier C, 15+ encounters), consistent with KG context
 340 being most valuable when document volume overwhelms chunk-level retrieval (Appendix L).

341 5.5 Physician Adjudication

342 A board-certified physician (A.S.) reviewed a stratified pilot sample ($n = 30$ ClinicalBench items,
 343 all 9 categories) under blinded conditions (condition labels randomized as A/B; unblinding per-
 344 formed post-review for analysis). Each item was rated on four dimensions: gold-standard correct-
 345 ness, model-answer correctness, clinical safety, and clinical utility. Full rubric definitions and the
 346 scoring interface are described in Appendix O.

347 **Physician-rated accuracy.** After unblinding, KG-RAG answers (C4g, $n = 16$) were rated correct
 348 in 81% of items vs. 29% for LLM-alone (C1, $n = 14$; Table 7), a +52 pp physician-rated delta—
 349 larger than the +47 pp auto-evaluator delta and directionally consistent.

350 **Safety and utility.** C4g answers were rated safe in 88% of cases (vs. 50% for C1), helpful in 81%
 351 (vs. 29%), and misleading in 6% (vs. 21%). No answers in either condition were rated potentially
 352 harmful.

Table 7: Physician adjudication pilot ($n = 30$). Single reviewer, blinded during scoring; condition labels unblinded post-review for stratified analysis.

Metric	C1 ($n = 14$)	C4g ($n = 16$)
Physician-rated correct	29%	81%
Partially correct	21%	6%
Incorrect	50%	12%
Safe	50%	88%
Minor concern	50%	12%
Potentially harmful	0%	0%
Helpful	29%	81%
Neutral	29%	12%
Not useful	21%	0%
Misleading	21%	6%

Auto-evaluator concordance. The keyword evaluator agreed with physician judgment in 47% of items. Disagreement was asymmetric: in 43% of cases the auto-evaluator was too strict (scored incorrect when the physician rated correct or partially correct), while in only 10% was it too lenient. This conservative bias is consistent with the abstention penalty noted in Section 5.1: C1 answers that correctly decline to answer without context are penalized by the keyword matcher.

Gold-standard quality. Physician review found 40% of auto-generated gold standards required revision (30% incorrect, 10% needing revision), with errors concentrated in change detection (0/4 correct) and negation (0/2 correct) categories. Gold-standard errors are constant across conditions and therefore do not affect the relative C1–C4g comparison, but they indicate that auto-generated clinical benchmarks require physician validation—a finding with implications beyond this study.

Limitations. The pilot is limited by small per-condition samples ($n = 14/16$) and a single reviewer; a second independent reviewer is planned. Despite these constraints, the physician-rated delta (+52 pp) is directionally and magnitude-consistent with the automated delta (+47 pp) across all four dimensions measured.

6 Discussion and Limitations

Why long context underperforms. Brute-force long context (C6) scores 39.0%, outperforming C1 (+17.2 pp) because it at least provides clinical text for the model to draw on, but it still falls far short of C4g (69.0%) and fails completely on historical (0%) and sequence (0%). C6 concatenates all notes chronologically without retrieval or selection; more sophisticated long-context strategies (hierarchical summarization, temporal segmentation) might perform better, but C6’s category-level failures—consistent with the “lost in the middle” phenomenon [37]—demonstrate that raw context dumping cannot substitute for structured retrieval. Historical questions require distinguishing resolved from active conditions across admissions; sequence questions require strict temporal ordering—both are implicit in raw text but explicit in the epistemic KG. C6 handles change detection modestly (23%) because the relevant information appears explicitly in the text, but KG-RAG achieves 100% (+77 pp).

Model-dependent KG benefit. Across the three models tested, KG-RAG benefit shows a consistent trend: models with lower C1 baselines gain more (Table 5). The two general-purpose models (Opus 21.8%, GPT-OSS 20.5%) gain +47 pp and +38 pp respectively, while the medically fine-tuned model (MedGemma, 26.2%) gains least (+32 pp). The pattern is consistent with a knowledge-compensator interpretation—structured epistemic context partially substitutes for parametric knowledge—though with only three models the trend is suggestive rather than conclusive. This has practical implications for RAG system design: the same retrieval pipeline may yield different gains depending on the downstream model, and single-model evaluation risks overstating or understating value.

When KG context matters. The complexity-dependent SliceBench effect (Tier C: +5.0 pp vs. Tier A: +0.6 pp) suggests structured epistemic context is most valuable when: (a) document volume

overwhelms chunk-level retrieval, (b) questions require cross-encounter synthesis where the same concept appears with different assertion states, and (c) assertion status evolves across encounters.

The C4 ablation cleanly decomposes the C3→C4g gap (+19.0 pp) into two components: assertions alone (C3→C4: −3.2 pp, CI crossing zero) and intent routing (C4→C4g: +22.2 pp, $p < 10^{-10}$). This reveals a key architectural insight: assertion metadata is a necessary but insufficient condition—its value is unlocked only when paired with question-type-specific retrieval. Without routing, assertions increase evidence verbosity and degrade temporal categories (historical: −30 pp, sequence: −47.5 pp); with routing, the same metadata enables near-perfect change detection (C4 10% → C4g 100%). The finding validates the two-component design and suggests that future work on assertion-aware clinical NLP should couple enriched metadata with structured retrieval strategies.

Uncertainty remains weakest. Uncertainty improves from 12.5% to 50.0% (+37.5 pp), but absolute accuracy remains low. Possible and suspected conditions are linguistically subtle (“consider statin,” “may have CHF”) and require pragmatic inference beyond keyword matching.

Threats to validity and claim scope. Our claims are bounded by dataset, evaluator, and model scope. Data scope is limited: ClinicalBench uses 43 patients (400 questions) and SliceBench 6 patients (144 questions), all from MIMIC-IV [33]; tier-level SliceBench effects remain hypothesis-generating ($n = 2$ patients per tier), and B2→B3 crosses zero (−1.5 pp, +5.9 pp). Evaluator uncertainty is substantial: keyword scoring agrees with physician judgment in 47% of the pilot, LLM-as-judge reaches 57%, and the three evaluators produce different absolute deltas. We therefore report all three and focus conclusions on directional agreement, not a single “true” effect size. Benchmark label quality is imperfect: physician review found a 40% error/revision rate in auto-generated gold answers (concentrated in change and negation), and pilot adjudication used a single reviewer. Model dependence is real: gains vary by model (Table 5), so Opus-scale effects should not be assumed universal. This study is an ablation analysis rather than a head-to-head leaderboard against existing medical RAG systems [5, 6, 15]; C2 is used as a proxy baseline for document-level retrieval on our longitudinal task. Runtime variance from quantized local inference (e.g., MedGemma ± 10 pp across runs) and single-site provenance further limit external validity. Finally, this work does *not* establish clinical deployment readiness: it evaluates QA behavior under controlled benchmarks, not prospective patient outcomes.

Conclusion. The epistemic propagation gap is a real and measurable phenomenon. EpiKG addresses it through end-to-end epistemic preservation combined with intent-aware retrieval that matches graph traversal to question semantics. The ablation ladder decomposes the gain into three independent contributions: structured retrieval (+28 pp, C1→C3), assertion metadata (−3 pp without routing, C3→C4), and intent-aware routing (+22 pp, C4→C4g). The counterintuitive C4 result—that assertions alone can degrade performance—shows that metadata enrichment without retrieval policy alignment is insufficient. Benefits are category-specific and model-dependent, and neither more context (C6) nor structure without reasoning (C7) closes the gap. The practical takeaway is a falsifiable design principle: preserve epistemic metadata *and* route retrieval by question intent, then evaluate by category×condition interaction rather than aggregate score alone.

References

- [1] Özlem Uzuner, Brett R. South, Shuying Shen, and Scott L. DuVall. 2010 i2b2/VA challenge on concepts, assertions, and relations in clinical text. *Journal of the American Medical Informatics Association*, 18(5):552–556, 2011.
- [2] Veysel Kocaman, Yigit Gul, M. Aytug Kaya, Hasham Ul Haq, Mehmet Butgul, Cabir Celik, and David Talby. Beyond negation detection: Comprehensive assertion detection models for clinical NLP. In *Text2Story Workshop at European Conference on Information Retrieval (ECIR)*, 2025.
- [3] OHDSI Collaborative. OMOP common data model v5.4. *Observational Health Data Sciences and Informatics*, 2024.
- [4] Darren Edge et al. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024.

- 441 [5] Tianjun Luo et al. GFM-RAG: Graph foundation model for retrieval augmented generation. In
442 *Advances in Neural Information Processing Systems*, 2025.
- 443 [6] Peng Jiang et al. KARE: Knowledge graph augmented reasoning via llms for clinical decision
444 support. In *International Conference on Learning Representations*, 2025.
- 445 [7] Junde Wu et al. Medical-Graph-RAG: Towards safe medical large language model via graph
446 retrieval-augmented generation. In *Proceedings of the 63rd Annual Meeting of the Association
447 for Computational Linguistics*, 2025.
- 448 [8] Hassan S. Al Khatib, Subash Neupane, Harish Kumar Manchukonda, Noorbakhsh Amiri Go-
449 lilarz, Sudip Mittal, Amin Amirlatifi, and Shahram Rahimi. Patient-centric knowledge graphs:
450 A survey of current methods, challenges, and applications. *Frontiers in Artificial Intelligence*,
451 7:1388479, 2024. doi: 10.3389/frai.2024.1388479.
- 452 [9] Bilal Abu-Salih, Muhammad Al-Qurishi, Mohammed Alweshah, Mohammad Al-Smadi,
453 Reem Alfayez, and Heba Saadeh. Healthcare knowledge graph construction: A systematic
454 review of the state-of-the-art, open issues, and opportunities. *Journal of Big Data*, 10(1):81,
455 2023. doi: 10.1186/s40537-023-00774-9.
- 456 [10] Marco Postiglione, Daniel Bean, Zeljko Kraljevic, Richard Dobson, and Vincenzo Moscato.
457 Predicting future disorders via temporal knowledge graphs and medical ontologies. *IEEE
458 Journal of Biomedical and Health Informatics*, 28(7):4238–4248, 2024. doi: 10.1109/JBHI.
459 2024.3390419.
- 460 [11] Daniel Ramirez et al. Graphiti: Building real-time knowledge graphs with temporal awareness.
461 *arXiv preprint arXiv:2501.13956*, 2025.
- 462 [12] James F. Allen. Maintaining knowledge about temporal intervals. *Communications of the
463 ACM*, 26(11):832–843, 1983.
- 464 [13] Fei Li, Jianfu Hong, Cui Tao, et al. TEO: A time event ontology for clinical narratives. *Journal
465 of the American Medical Informatics Association*, 27(10):1560–1568, 2020.
- 466 [14] Yan Huang, Xiaojin Li, and Guo-Qiang Zhang. Temporal cohort logic. *AMIA Annual Sympo-
467 sium Proceedings*, 2022:1237–1246, 2023.
- 468 [15] Yifan Lu, Tianyu Fu, et al. Doctorrag: Medical rag emulating doctor-like reasoning. In *Ad-
469 vances in Neural Information Processing Systems*, 2025.
- 470 [16] Yusheng Wang et al. MedRAG: Enhancing medical diagnosis through retrieval-augmented
471 generation with knowledge graph-elicited reasoning. *Proceedings of The Web Conference*,
472 2025.
- 473 [17] Lang Cao, Qingyu Chen, and Yue Guo. EHR-RAG: Bridging long-horizon structured elec-
474 tronic health records and large language models via enhanced retrieval-augmented generation.
475 *arXiv preprint arXiv:2601.21340*, 2026.
- 476 [18] Samuel Thio, Matthew Lewis, Spiros Denaxas, and Richard J. B. Dobson. Unlocking elec-
477 tronic health records: A hybrid graph RAG approach to safe clinical AI for patient QA. *arXiv
478 preprint arXiv:2602.00009*, 2025.
- 479 [19] Boci Peng, Yun Zhu, Yongchao Liu, Xiaohe Bo, Haizhou Shi, Chuntao Hong, Yan Zhang,
480 and Siliang Tang. Graph retrieval-augmented generation: A survey. *ACM Transactions on
481 Information Systems*, 2025. doi: 10.1145/3777378.
- 482 [20] Wei Chen et al. Multi-LLM KG-RAG: End-to-end clinical knowledge graph construction.
483 *arXiv preprint arXiv:2601.01844*, 2026.
- 484 [21] Lang Li et al. AutoRD: An automatic and end-to-end system for rare disease knowledge graph
485 construction. *JMIR Medical Informatics*, 12, 2024.
- 486 [22] Rakhilya Lee Mekhtieva, Brandon Forbes, Dalal Alrajeh, Brendan Delaney, and Alessandra
487 Russo. RECAP-KG: Mining knowledge graphs from raw GP notes for remote COVID-19
488 assessment in primary care. *arXiv preprint arXiv:2306.17175*, 2023.
- 489 [23] Guohui Xiao, Emily Pfaff, Eric Prud’hommeaux, David Booth, Deepak K. Sharma, Nan Huo,
490 Yue Yu, Nansu Zong, Kathryn J. Ruddy, Christopher G. Chute, and Guoqian Jiang. FHIR-
491 Ontop-OMOP: Building clinical knowledge graphs in FHIR RDF with the OMOP common

- data model. *Journal of Biomedical Informatics*, 134:104201, 2022. doi: 10.1016/j.jbi.2022.104201.
- [24] Wendy W. Chapman, Will Bridewell, Paul Hanbury, Gregory F. Cooper, and Bruce G. Buchanan. A simple algorithm for identifying negated findings and diseases in discharge summaries. *Journal of Biomedical Informatics*, 34(5):301–310, 2001.
- [25] Henk Harkema, John N. Dowling, Tyler Thornblade, and Wendy W. Chapman. ConText: An algorithm for determining negation, experience, and temporal status from clinical reports. *Journal of Biomedical Informatics*, 42(5):839–851, 2009.
- [26] Richard T. Snodgrass. *Developing Time-Oriented Database Applications in SQL*. Morgan Kaufmann, 2000.
- [27] Adil Bahaj and Mounir Ghogho. A step towards quantifying, modelling and exploring uncertainty in biomedical knowledge graphs. *Computers in Biology and Medicine*, 184:109355, 2025. doi: 10.1016/j.combiomed.2024.109355.
- [28] Yifan Peng, Xiaosong Wang, Le Lu, Mohammadhadi Bagheri, Ronald Summers, and Zhiyong Lu. NegBio: A high-performance tool for negation and uncertainty detection in radiology reports. *AMIA Summits on Translational Science Proceedings*, 2018.
- [29] Werner Ceusters and Barry Smith. Aboutness: Towards foundations for the information artifact ontology. *International Conference on Biomedical Ontology*, 2015.
- [30] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948.
- [31] Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *Applied Sciences*, 11(14):6421, 2021. doi: 10.3390/app11146421.
- [32] Yixing Jiang, Kameron C. Black, Gloria Geng, Danny Park, James Zou, Andrew Y. Ng, and Jonathan H. Chen. MedAgentBench: A virtual EHR environment to benchmark medical LLM agents. *NEJM AI*, 2(9), 2025. doi: 10.1056/AIdbp2500144.
- [33] Alistair E.W. Johnson, Lucas Bulgarelli, Lu Shen, et al. MIMIC-IV, a freely accessible electronic health record dataset. *Scientific Data*, 10:1, 2023.
- [34] Guangzhi Xiong et al. MIRAGE: Medical information retrieval-augmented generation evaluation. In *Findings of the Association for Computational Linguistics: ACL*, 2024.
- [35] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993.
- [36] Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
- [37] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.
- [38] Yanjun Gao, Ruizhe Li, Emma Croxford, John Caskey, Brian W. Patterson, Matthew Churpek, Timothy Miller, Dmitriy Dligach, and Majid Afshar. Leveraging medical knowledge graphs into large language models for diagnosis prediction: Design and application study. *JMIR AI*, 4(1):e58670, 2025. doi: 10.2196/58670.
- [39] Gautier Izacard, Mathilde Caron, Lucas Hosseini, Sebastian Riedel, Piotr Bojanowski, Armand Joulin, and Edouard Grave. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research*, 2022.

536 A Notation Summary

Symbol	Meaning
$\alpha \in \mathcal{A}$	Assertion label (7 values, Eq. 1)
τ_v	Valid time (event date, valid from/to)
τ_t	Transaction time (recorded at, doc date, created at)
τ_a	NLP-asserted temporality (Current / Past / Future)
$r \in \mathcal{R}$	Temporal relation (9 values, Table 20)
c	Confidence score $\in [0, 1]$
ξ	Experiencer (Patient / Family)
$e(m)$	Epistemic state tuple (c, α, ξ, τ)
q, π	Clinical question, patient
ι	Question intent (Change / Current_State / Historical / Default)
$f_{np}(c)$	Fraction of non-present assertions for concept c

Table 8: Key notation used throughout the paper.

537 B Formal Epistemic Preservation

538 We formalize the epistemic invariant maintained by the EpiKG pipeline, building on the principle
 539 that aboutness—the relationship between information artifacts and the entities they represent—must
 540 be preserved across transformations [29].

541 **Definition 2** (Epistemic State). *For a clinical mention m , the epistemic state is the tuple $e(m) =$
 542 (c, α, ξ, τ) , where c is the OMOP concept identifier, $\alpha \in \mathcal{A}$ is the 7-value assertion label (Eq. 1), $\xi \in$
 543 $\{\text{PATIENT, FAMILY}\}$ is the experiencer, and $\tau \in \{\text{CURRENT, PAST, FUTURE}\}$ is the temporality. A
 544 pipeline P epistemically preserves mention m if $P(e(m)) = e(m)$; that is, the epistemic state at the
 545 output of every pipeline stage is identical to the state at extraction.*

Proposition 3 (Assertion Entropy Loss). *For concept c in patient π ’s record, let A_c^π denote the
 assertion distribution with empirical frequencies $p(\alpha_i) = n_i / \sum_j n_j$ over the $|\mathcal{A}|$ assertion classes.
 The assertion entropy is*

$$H(A_c^\pi) = - \sum_i p(\alpha_i) \log p(\alpha_i) . \quad (3)$$

546 *An assertion-blind pipeline collapses all mentions to $\alpha = \text{PRESENT}$, yielding a degenerate distribu-*
 547 *tion with $H = 0$. The information loss is $\Delta H = H(A_c^\pi) \geq 0$ [30], with $\Delta H > 0$ strictly whenever*
 548 *any mention carries a non-present assertion.*

549 *Proof.* Under the assertion-blind mapping $\phi : \alpha_i \mapsto \text{PRESENT}$ for all i , the output distribution
 550 assigns probability 1 to PRESENT and 0 to all other classes, so $H(\phi(A_c^\pi)) = 0$. By non-negativity
 551 of Shannon entropy, $\Delta H = H(A_c^\pi) - 0 = H(A_c^\pi) \geq 0$, with equality iff all mentions already carry
 552 $\alpha = \text{PRESENT}$. $\square$

553 **Corollary 4** (Faithfulness Bound). *Let $f_{np}(c)$ denote the fraction of mentions of concept c carrying*
 554 *a non-present assertion ($\alpha \neq \text{PRESENT}$). Without assertion labels, the maximum assertion-faithful*
 555 *accuracy for any downstream task conditioned on c is bounded by $1 - f_{np}(c)$. In clinical records*
 556 *where negated and uncertain mentions are prevalent—e.g., “no pneumonia,” “possible CHF”—this*
 557 *bound can be substantially below 1.*

558 *Proof.* Without assertion labels, any predictor must assign a single assertion class to all mentions
 559 of c . Choosing PRESENT (the majority class in clinical text) yields accuracy $1 - f_{np}(c)$; the $f_{np}(c)$
 560 non-present mentions are necessarily misclassified. $\square$

561 C Gap Analysis

562 Partial marks indicate limited coverage: GFM-RAG, MedRAG, and KARE build population-level
 563 or hierarchical KGs (not per-patient longitudinal graphs); KARE uses KG-augmented retrieval but

Table 9: Capability comparison across representative systems. ✓ = supported, ○ = partial, ✗ = not supported.

System	Patient KG	OMOP mapping	Assertion (7-class)	Tri-temporal	Assert-aware RAG	Experienter	Allen's algebra	Multi-hop (≥ 3)
DoctorRAG [15]	✗	✗	✗	✗	✗	✗	✗	✗
GFM-RAG [5]	○	✗	✗	✗	✗	✗	✗	✓
MedRAG [16]	○	✗	✗	✗	✗	✗	✗	○
KARE [6]	○	✗	✗	✗	○	✗	✗	✓
Multi-LLM KG [20]	✗	✓	○	✗	✗	✗	✗	✗
MedTKG [10]	✓	✗	✗	○	✗	✗	✗	✓
Graphiti [11]	✗	✗	✗	○	✗	✗	✗	✓
MediGRAF [18]	✓	✗	✗	✗	✗	✗	✗	✓
EpiKG (ours)	✓	✓	✓	✓	✓	✓	✓	○

without assertion awareness; Multi-LLM KG models uncertainty via entropy but not the full assertion taxonomy; MedTKG and Graphiti use temporal annotations but not tri-temporal models. MediGRAF is the closest patient-level competitor, constructing per-patient medical graphs from clinical notes, but it does not preserve assertion status or temporal relations as edge attributes. EpiKG’s partial mark on multi-hop traversal (≥ 3 hops) reflects a deliberate architectural trade-off: PostgreSQL CTE-based traversal provides ACID compliance but degrades beyond 2 hops (Appendix K).

D Safety Score

The safety score $S = 1 - \frac{1}{N} \sum_{i=1}^N w_i \cdot \mathbb{I}[\hat{y}_i \neq y_i]$ applies $w_i = 2.0$ for false-positive assertion errors (reporting a negated or absent condition as present) and $w_i = 1.0$ otherwise, capturing asymmetric clinical risk where acting on a falsely affirmed condition is more dangerous than missing a present one. We treat $w = 2.0$ as a reasonable default; sensitivity to this choice is a direction for future work.

E Cross-Model Per-Category Results

Table 10: ClinicalBench per-category accuracy (%) by model under C1 and C4g (intent-aware KG-RAG, evaluator v2). All three models improve on all 9 categories.

Category	n	Claude Opus 4.6			MedGemma 27B			GPT-OSS 20B		
		C1	C4g	Δ	C1	C4g	Δ	C1	C4g	Δ
Negation	110	45.5	80.9	↑35.5 pp	57.3	76.4	↑19.1 pp	46.4	80.0	↑33.6 pp
Conditional	20	0.0	45.0	↑45.0 pp	15.0	35.0	↑20.0 pp	0.0	15.0	↑15.0 pp
Uncertainty	40	12.5	50.0	↑37.5 pp	7.5	32.5	↑25.0 pp	5.0	30.0	↑25.0 pp
Family hist.	30	3.3	56.7	↑53.3 pp	3.3	40.0	↑36.7 pp	3.3	40.0	↑36.7 pp
Sequence	40	10.0	65.0	↑55.0 pp	7.5	47.5	↑40.0 pp	12.5	42.5	↑30.0 pp
Current st.	50	26.0	70.0	↑44.0 pp	32.0	58.0	↑26.0 pp	36.0	56.0	↑20.0 pp
Duration	30	30.0	66.7	↑36.7 pp	53.3	70.0	↑16.7 pp	16.7	70.0	↑53.3 pp
Historical	50	6.0	60.0	↑54.0 pp	0.0	46.0	↑46.0 pp	0.0	60.0	↑60.0 pp
Change	30	6.7	100.0	↑93.3 pp	0.0	76.7	↑76.7 pp	0.0	70.0	↑70.0 pp
Overall	400	21.8	69.0	↑47.2 pp	26.2	57.8	↑31.5 pp	20.5	58.0	↑37.5 pp

F MedGemma Full Per-Category Results

Category	n	C7	C1	C4g
Negation	110	13.6	57.3	76.4
Conditional	20	0.0	15.0	35.0
Uncertainty	40	0.0	7.5	32.5
Family History	30	0.0	3.3	40.0
Sequence	40	0.0	7.5	47.5
Current State	50	0.0	32.0	58.0
Duration	30	0.0	53.3	70.0
Historical	50	0.0	0.0	46.0
Change	30	0.0	0.0	76.7
Overall	400	3.8	26.2	57.8

Table 11: MedGemma 27B ClinicalBench per-category accuracy (%; keyword evaluator v2) for three conditions. MedGemma gains +32 pp overall under C4g, with all 9 categories improving. Strongest improvements on change (+77 pp) and historical (+46 pp).

578 G Experienter Attribute Propagation Ablation

579 The experienter attribute distinguishes patient conditions from family member conditions;
580 without it, the graph conflates the two, causing family history misattribution.

Table 12: Impact of experienter attribute propagation on Opus C4 (400 questions, prior evaluator run). Categories with no change omitted.

Category	Pre-fix	Post-fix	Δ
Family history	63.3%	73.3%	+10.0 pp
Uncertainty	50.0%	55.0%	+5.0 pp
Historical	34.0%	38.0%	+4.0 pp
Current state	46.0%	48.0%	+2.0 pp
Overall C4	68.2%	70.2%	+2.0 pp

581 The fix improved family history by +10.0 pp with zero regressions on guard categories (negation,
582 conditional, duration, sequence all unchanged), confirming that the experienter attribute is load-
583 bearing for assertion-sensitive reasoning.

584 H SliceBench Figures




Figure 3: SliceBench results ($n = 2$ patients per tier, 24 questions each; point estimates without per-tier CIs). **Top:** The KG layer’s incremental benefit (B2→B3) shows a directional, complexity-dependent trend: +0.6 pp (Tier A) to +5.0 pp (Tier C). **Bottom:** Full condition progression by tier. The overall B2→B3 delta is +2.2 pp (CI: [-1.5 pp, +5.9 pp]).

I Per-Category Delta Chart




Figure 4: Per-category accuracy change (Claude Opus 4.6, evaluator v2) decomposing the C3→C4g gain. Categories sorted by C3→C4 delta; dashed line separates Task A (assertion-sensitive) from Task B (temporal). **Left:** C3→C4 (assertions without routing): assertion categories improve (negation +24 pp, conditional +20 pp) but temporal categories collapse (sequence −48 pp, historical −30 pp). **Middle:** C4→C4g (intent routing): routing recovers temporal categories (change +90 pp, sequence +55 pp, historical +48 pp) while slightly reducing assertion categories. **Right:** C1→C4g (overall): all 9 categories improve.

586 J SliceBench Conditions

Table 13: SliceBench conditions. B0–B4 form a monotone context progression.

ID	Condition	Context Provided to LLM
B0	LLM Alone	Question only (parametric knowledge)
B1	Latest Note	Most recent clinical note
B2	All Notes RAG	All notes, retrieved by relevance
B3	KG-RAG	Knowledge graph context + all notes
B4	Full System	KG-RAG + guidelines + calculators

587 K System Implementation Details

588 **LLM baseline.** On MedQA-USMLE (965 questions), Claude Opus 4.5 achieves 81.6% accu-
589 racy (5.1 pp below GPT-4 at 86.7% and Med-PaLM 2 at 86.5% [32]), establishing that the Claude
590 model family—Sonnet 4.5 (SliceBench answerer) and Opus 4.6 (ClinicalBench primary, SliceBench
591 judge)—provides a reasonable LLM baseline (Table 14).

Model	Accuracy	Step 1	N
EpiKG (LLM alone)	81.6%	79.9%	965
GPT-4 (2023)	86.7%	—	1,273
Med-PaLM 2 (2023)	86.5%	—	1,273
Claude 3 Opus (2024)	78.0%	—	1,273

Table 14: MedQA-USMLE results. Our LLM-alone baseline is competitive.

592 **KG scale and traversal latency.** The deployed system processes 145 documents across 85 pa-
593 tients, materializing 3,100 KG nodes and 8,803 edges. Graph traversal latency is sub-millisecond
594 at this scale: 0.57 ms (1-hop), 0.75 ms (2-hop). The full system integrates 201 clinical calculators,
595 1,202 guideline sections, 20M+ OMOP vocabulary relationships, 122 assertion trigger patterns, and
596 24 edge types across 13 node types.

597 **Multi-hop traversal.** On the DR.KNOWS diagnostic reasoning benchmark [38], which measures
598 KG traversal accuracy using PostgreSQL-backed clinical data, EpiKG achieves 0.420 overall (50%

at 1-hop, 25% at 2-hop, 0% at 3-hop). Multi-hop degradation reflects a deliberate architectural trade-off: PostgreSQL CTE-based traversal provides ACID compliance but scales poorly beyond 2 hops.

L SliceBench Tier-Stratified Results

Table 15: SliceBench results stratified by patient complexity tier. B2→B3 Δ shows the incremental KG contribution.

Tier	Encounters	Score (%)					B2→B3 Δ
		B0	B1	B2	B3	B4	
A	1–2 notes	51.7	66.7	81.1	81.8	82.6	+0.6 pp
B	5–10 notes	49.5	79.4	86.4	87.4	90.1	+1.0 pp
C	15+ notes	48.4	74.2	86.5	91.5	90.3	+5.0 pp

M ClinicalBench Full Per-Category Results (Opus)

Category	n	C7	C6	C1	C2	C2b	C3	C4	C4g
Negation	110	13.6	69.1	45.5	70.9	75.5	65.5	90.0	80.9
Conditional	20	0.0	35.0	0.0	35.0	50.0	30.0	50.0	45.0
Uncertainty	40	0.0	27.5	12.5	35.0	32.5	37.5	52.5	50.0
Family History	30	0.0	43.3	3.3	43.3	43.3	43.3	50.0	56.7
Sequence	40	0.0	0.0	10.0	60.0	47.5	57.5	10.0	65.0
Current State	50	0.0	44.0	26.0	32.0	32.0	54.0	30.0	70.0
Duration	30	0.0	66.7	30.0	66.7	66.7	60.0	46.7	66.7
Historical	50	0.0	0.0	6.0	48.0	46.0	42.0	12.0	60.0
Change	30	0.0	23.3	6.7	43.3	20.0	16.7	10.0	100.0
Overall	400	3.8	39.0	21.8	52.2	50.7	50.0	46.8	69.0

Table 16: Complete ClinicalBench per-category accuracy (%), Claude Opus 4.6, keyword evaluator v2) for all conditions. C2b (Contriever dense retrieval) scores comparably to C2 (TF-IDF), confirming the retrieval method does not drive the C2→C4g gap. C4 adds assertion metadata to C3 without intent routing; it helps assertion-sensitive categories (negation: +24.5 pp) but hurts temporal categories (historical: −30 pp, sequence: −47.5 pp). Only with intent routing (C4g) does the full system reach 69%.

N Physician Validation Pilot

A board-certified physician (A.S.) reviewed a stratified sample of $n = 30$ ClinicalBench items (14 C1, 16 C4g) spanning all 9 categories under blinded conditions. Each item was rated on: (1) gold-standard correctness (Yes / Partially / Needs Revision / No), (2) model-answer correctness (Correct / Partially Correct / Incorrect), (3) clinical safety (Safe / Minor Concern / Potentially Harmful), and (4) clinical utility (Helpful / Neutral / Not Useful / Misleading). Reviews were conducted using a custom scoring interface with access to the full de-identified MIMIC-IV discharge summaries, filtered to the admission(s) relevant to each question.

Gold-standard error analysis. Of the 30 gold standards reviewed, 18 (60%) were rated correct, 3 (10%) needed revision, and 9 (30%) were incorrect. Errors were concentrated in change detection (0/4 gold standards correct—NLP misread discharge medication lists) and negation (0/2 correct—NLP missed physical exam findings). These errors are constant across conditions and do not affect the relative C1–C4g comparison.

Auto-evaluator concordance. The keyword evaluator agreed with physician judgment in 47% of cases. In 43% of disagreements the evaluator was too strict (marked incorrect when the physician rated correct or partially correct); in 10% it was too lenient. This conservative bias largely reflects the abstention penalty: C1 answers that appropriately decline to answer without patient context are penalized by keyword matching.

O Physician Validation Protocol

Design. A board-certified physician independently scored ClinicalBench answers under two conditions (C1 and C4g) using a multi-dimensional rubric (correctness, safety, utility). The reviewer was blinded to condition assignment (conditions randomized as A/B). The pilot sample comprises 30 items stratified across all 9 categories and both conditions, selected to maximize coverage within a feasible review budget (~8 hours physician time).

Endpoints. (1) Human-keyword evaluator concordance rate (fraction of automated scores confirmed by physician). (2) Physician-rated C4g vs. C1 accuracy, safety, and utility. (3) Gold-standard error rate.

Disclosure. The reviewing physician (A.S.) is a co-author. Single-reviewer design limits inter-rater reliability assessment; a second reviewer is planned for the extended validation.

Status. This protocol was designed before examining primary results. Pilot results ($n = 30$) are reported in Section 5.5.

P Reproducibility Details

Models. ClinicalBench primary ablation: Claude Opus 4.6 (claude-opus-4-20250514, keyword evaluator v2). Cross-model: MedGemma 27B (gemma3:27b, 4-bit GGUF) and GPT-OSS 20B (4-bit GGUF) via Ollama; same evaluator. SliceBench: Claude Sonnet 4.5 (claude-sonnet-4-5-20250929) answers, Opus 4.6 judges. Temperature 0 throughout; 4-bit quantization introduces GPU non-determinism (± 10 pp run-to-run for MedGemma), controlled via within-run paired comparisons.

Evaluation. The deterministic keyword evaluator performs exact word-boundary matching (regex `\bKEYWORD\b`) against gold-standard assertion and temporal keywords. Per-category keyword sets: negation (“no,” “denies,” “absent,” “negative”), uncertainty (“possible,” “suspected,” “may,” “likely,” “consider”), temporal (“before,” “after,” “during,” “changed,” “new”), plus category-specific terms. Evidence preambles (echoed graph context) are stripped before matching.

Retrieval. C2’s vanilla RAG uses TF-IDF over chunked clinical notes (512-token chunks, 64-token overlap), retrieving the top-5 chunks by cosine similarity. C2b replaces TF-IDF with Contriever [39] dense embeddings (256-token chunks, 64-token overlap, top- k up to 6,000 characters); C2b scores 50.7% ($\Delta = -1.5$ pp vs. C2; $p = 0.62$), confirming the retrieval method does not explain the C2→C4g gap. These are the closest analogs to existing medical RAG systems (Doctor-RAG, MedRAG) on patient-level longitudinal data. ClinicalBench conditions map to SliceBench: C1≈B0, C2≈B2, C4g≈B3, C5≈B4.

Bootstrap. $n = 2,000$ resamples, seed 42, BCa method [35]. We report both question-level and patient-level (cluster) bootstrap CIs; the latter resample the 43 patients with replacement, including all their questions per draw, to account for within-patient correlation. Patient-level CIs are slightly wider but all primary endpoints remain significant (Table 17).

McNemar’s test. We supplement bootstrap CIs with McNemar’s test for paired nominal data, comparing discordant pairs (C1 wrong/C4g right vs. C1 right/C4g wrong). All three pairwise comparisons are significant: C1 vs. C4g ($\chi^2 = 157.1$, $p < 10^{-6}$; discordant: 18 vs. 207), C1 vs. C3 ($\chi^2 = 73.4$, $p < 10^{-6}$; 29 vs. 142), and C3 vs. C4g ($\chi^2 = 49.3$, $p < 10^{-11}$; 19 vs. 95).

Data. De-identified MIMIC-IV clinical notes accessed under a PhysioNet Credentialed Health Data Use Agreement (v1.5.0).

Compute. ClinicalBench Opus C1/C3/C4g: ~22min each (API); Opus C6: ~3.5h (API); C7: < 1min (no LLM); MedGemma C1–C4g: ~3.6h (single GPU); MedGemma C6: ~1.5h; SliceBench: ~2h (API).

Evaluator evolution. The keyword evaluator underwent three versions: v0 (substring matching), v1 (word-boundary matching + evidence preamble stripping), and v2 (+ abstention detection gate + domain-specific keyword requirements for sequence and change). The v1 evaluator awarded false positives when models responded with “insufficient information in the notes” because negation and

Table 17: Statistical summary: question-level and patient-level (cluster) BCa bootstrap 95% CIs for primary endpoints. Both levels exclude zero for all comparisons.

Comparison	Δ	Question CI	Patient CI	McNemar p
C4g—C1 (overall)	+47.2 pp	[+40.9 pp, +52.8 pp]	[+40.2 pp, +52.5 pp]	$< 10^{-6}$
C2b—C2 (dense vs. keyword)	−1.5 pp	[−6.6 pp, +3.5 pp]	[−5.9 pp, +3.0 pp]	0.62
C4g—C2b (KG-RAG vs. dense)	+18.2 pp	[+12.3 pp, +23.5 pp]	[+12.9 pp, +23.2 pp]	$< 10^{-9}$
C3—C1 (retrieval)	+28.2 pp	[+22.3 pp, +34.0 pp]	[+23.1 pp, +33.8 pp]	$< 10^{-6}$
C4—C3 (assertions)	−3.2 pp	[−10.0 pp, +3.2 pp]	[−10.8 pp, +3.4 pp]	0.37
C4g—C4 (routing)	+22.2 pp	[+16.0 pp, +28.0 pp]	[+14.8 pp, +29.3 pp]	$< 10^{-10}$
C4g—C3 (both)	+19.0 pp	[+13.8 pp, +24.0 pp]	[+14.7 pp, +23.5 pp]	$< 10^{-11}$
Hard longitudinal	+59.2 pp	[+49.2 pp, +67.7 pp]	[+48.6 pp, +69.4 pp]	—

temporal keywords in the refusal text matched gold-standard patterns. The v2 evaluator adds an abstention detection layer: answers matching abstention patterns (e.g., “cannot determine,” “not mentioned,” “insufficient evidence”) are scored as incorrect unless they contain clinical claim patterns (e.g., “patient does not,” “denies”). Additionally, v2 requires sequence answers to contain ordering keywords (“first,” “then,” “before”) and change answers to contain change keywords (“added,” “removed,” “discontinued”), preventing term-overlap-only false positives. The duration category’s minimum-0.5 score floor for matching duration keywords was removed. The v1→v2 transition reduced C1 accuracy (from ~50% to 21.8% for Opus) by correctly classifying abstention responses as incorrect; C4g accuracy decreased modestly (from ~76% to 69.0%) because the model abstains less frequently when retrieval context is provided. All main-text numbers use v2.

Reproducibility package. A standalone reproducibility package is included as supplementary material (epikg-benchmark/). It contains all 400 ClinicalBench questions with gold answers, raw model predictions for all LLM-based conditions (Opus C1/C2/C3/C4/C4g/C6, MedGemma C1/C4g, GPT-OSS C1/C4g), scored outputs for the deterministic baseline (C7), and the keyword evaluator v2 with abstention detection. Running the evaluator requires only Python 3.10+ with no external dependencies. All accuracy numbers in Tables 3, 4, and 5 are independently reproducible from this package. MIMIC-IV patient identifiers are included so that PhysioNet-credentialed reviewers can trace predictions back to source clinical notes.

Results provenance. Table 18 maps each reported result to its source checkpoint file. MedGemma C4g has 399 predictions (1 timeout); the missing question is scored as incorrect ($n = 400$).

Table 18: Results provenance: checkpoint files for each condition×model combination. All scored with keyword evaluator v2.

Condition	Model	Checkpoint file	n
C1	Opus 4.6	opus/C1_llm_alone.json	400
C2	Opus 4.6	opus/C2_vanilla_rag.json	400
C2b	Opus 4.6	opus/C2b_dense_rag.json	400
C3	Opus 4.6	opus/C3_kg_rag.json	400
C4	Opus 4.6	opus/C4_epistemic_kg_rag.json	400
C4g	Opus 4.6	opus/C4g_intent_aware.json	400
C6	Opus 4.6	opus/C6_long_context.json	400
C7	—	opus/C7_deterministic.json	400
C1	MedGemma 27B	medgemma/C1_llm_alone.json	400
C4g	MedGemma 27B	medgemma/C4g_intent_aware.json	399*
C1	GPT-OSS 20B	gptoss/C1_llm_alone.json	400
C4g	GPT-OSS 20B	gptoss/C4g_intent_aware.json	400

*MedGemma C4g: 1 timeout; missing answer scored as incorrect for $n = 400$ denominator.

Q Assertion Category Definitions

Assertion	Definition	Example
PRESENT	Condition currently affirmed	“has diabetes”
ABSENT	Condition explicitly negated	“denies chest pain”
POSSIBLE	Condition suspected, not confirmed	“possible pneumonia”
CONDITIONAL	Contingent on specific circumstances	“if febrile, start antibiotics”
HYPOTHETICAL	Discussed as a scenario	“would need dialysis if. . .”
FAMILY_HISTORY	Attributed to a family member	“mother had breast cancer”
HISTORICAL	Previously true, not necessarily current	“former smoker”

Table 19: The 7-value assertion taxonomy, extending i2b2 by separating HISTORICAL from FAMILY_HISTORY.

R Temporal Relation Mapping

The nine temporal relations $\mathcal{R}$ used on KG edges are derived from Allen’s 13 canonical interval relations [12] by merging symmetric pairs.

$\mathcal{R}$ value	Allen source(s)	Meaning
BEFORE	Before, Meets	A ends before B starts
AFTER	After, Met-by	A starts after B ends
DURING	During	A entirely within B
CONTAINS	Contains	A entirely contains B
OVERLAPS	Overlaps, Overlapped-by	A and B partially overlap
STARTS	Starts, Started-by	A and B share start time
FINISHES	Finishes, Finished-by	A and B share end time
CONCURRENT	Equals	A and B have same interval
UNKNOWN	—	Relation undetermined

Table 20: Mapping from Allen’s 13 interval relations to the 9 temporal relation values stored on KG edges.

S Intent-Aware Routing Algorithm

The C4g classifier is rule-based, mapping questions to one of four intents using keyword patterns (e.g., “changed”, “new since” trigger CHANGE; “currently”, “active problem” trigger CURRENT_STATE). Algorithm 1 details the full retrieval procedure.

Algorithm 1: Intent-Aware Retrieval (C4g)

Input: Question q , patient π
Output: Structured evidence E

```

1  $\mathcal{C} \leftarrow \text{EXTRACTCONCEPTS}(q)$ ; // NLP + OMOP enrichment
2  $\iota \leftarrow \text{CLASSIFYINTENT}(q)$ ; //  $\in \{\text{CHANGE, CURRST, HIST, DEFAULT}\}$ 
3 if  $\iota = \text{CHANGE}$  then
4   Partition edges by admission:  $\mathcal{E}_k \leftarrow \{e \mid \text{hadm\_id}(e) = k\}$ ;
5   foreach admission pair  $(k, k')$  with  $k < k'$  do
6      $\mathcal{A} \leftarrow \mathcal{C}_{k'} \setminus \mathcal{C}_k$ ;  $\mathcal{R} \leftarrow \mathcal{C}_k \setminus \mathcal{C}_{k'}$ ;  $\mathcal{S} \leftarrow \mathcal{C}_k \cap \mathcal{C}_{k'}$ ;
7      $E_g \leftarrow \text{FORMATCHANGE}(\mathcal{A}, \mathcal{R}, \mathcal{S})$ ;
8 else if  $\iota = \text{CURRST}$  then
9    $E_g \leftarrow \text{FILTEREDGES}(\pi, \mathcal{C}, \tau_a = \text{CURRENT} \vee \text{open validity})$ ;
10  Deduplicate by concept; emit “NOT FOUND” for missing  $c \in \mathcal{C}$ ;
11 else if  $\iota = \text{HIST}$  then
12   $E_g \leftarrow \text{FILTEREDGES}(\pi, \mathcal{C}, \tau_a = \text{PAST})$ ;
13  Augment: concepts in earlier admissions but absent from latest  $\rightarrow$  “resolved”;
14 else
15   $E_g \leftarrow \text{BIDIRECTIONALBFS}(\pi, \mathcal{C}, \text{hops} = 2\text{--}3, c_{\min} = 0.3)$ ;
16  $E_d \leftarrow \text{RETRIEVEDOCUMENTS}(\pi, \mathcal{C})$ ;
17 return  $\text{COMPOSE}(E_g, E_d, \text{template}(\iota))$ ;

```

T Bookend and Diagnostic Conditions

C5 (Full System). C5 extends C4 with guideline retrieval (1,202 sections) and clinical calculators (201, e.g., CHA₂DS₂-VASc, MELD). In a prior evaluator run, C5 achieved 68.2% overall, below the prior-run C1 and C4, likely because non-relevant components dilute the context window.

C6 (Long Context). All patient documents are concatenated chronologically and presented to the LLM. For Opus (200K token window), all notes fit; for MedGemma (8K window), later documents are truncated. C6 achieves 39.0% (Opus) overall—better than C1 (21.8%) but far below C4g (69.0%). Opus scores 0% on historical and sequence under C6, confirming that brute-force long context cannot substitute for structured retrieval on categories requiring cross-encounter synthesis.

C7 (Deterministic KG). KG edges matching query concepts are returned directly without LLM reasoning. C7 achieves 3.8% overall (only negation at 13.6% via stored assertion metadata), confirming that structured data without an LLM reasoner cannot answer clinical questions.

U Illustrative Retrieval Example

To illustrate why intent-aware routing matters, consider a change question: “*What medications changed between this patient’s first and second admissions?*”

C1 (LLM alone). The model sees only the current note, which mentions metoprolol, atorvastatin, and “discontinued lisinopril.” Without prior admission records, it cannot determine what was *added* vs. *continued*, and may hallucinate prior medication lists.

C4g (intent-aware KG-RAG). The intent classifier routes to CHANGE retrieval, which partitions KG edges by `hadm_id` and computes set differences:

- **Added** (admission 2 only): atorvastatin 40mg [PRESENT]
- **Removed** (admission 1 only): lisinopril 10mg [PRESENT→ABSENT]
- **Continued**: metoprolol 25mg [PRESENT, both admissions]

The assertion labels (PRESENT, ABSENT) disambiguate: “discontinued lisinopril” is not a current medication but a historical one whose status changed—exactly the distinction the epistemic schema preserves.

V LLM-as-Judge Concordance

To complement the deterministic keyword evaluator, we rescored all 800 predictions (400 questions $\times$ 2 conditions: C1 and C4g) using Claude Opus 4.6 as an LLM judge. The judge prompt presents the question, reference answer, and system answer, and requests a score of 1 (correct), 0.5 (partially correct), or 0 (incorrect) with a one-sentence justification. Table 21 summarizes the concordance.

Table 21: Evaluator concordance comparison. LLM judge scores ≥ 0.5 treated as correct for binary comparison. Physician concordance computed on $n = 30$ overlapping items.

Metric	Keyword v2	LLM Judge
<i>vs. Keyword evaluator ($n = 800$)</i>		
Agreement	—	78.9%
Cohen’s κ	—	0.572
<i>vs. Physician ($n = 30$)</i>		
Agreement	46.7%	56.7%
Too strict	43.3%	36.7%
Too lenient	10.0%	6.7%
<i>Condition-level accuracy</i>		
C1 accuracy	21.8%	28.5%
C4g accuracy	69.0%	57.0%
C4g–C1 Δ	+47.2 pp	+28.5 pp

Key findings. The LLM judge achieves higher physician concordance (56.7% vs. 46.7%) and is less prone to false strictness (36.7% vs. 43.3%), confirming that the keyword evaluator’s conservative bias inflates the measured delta. Under the LLM judge, C1 rises from 21.8% to 28.5% (the judge credits clinically reasonable hedging that keyword matching penalizes), while C4g drops from 69.0% to 57.0% (the judge penalizes partial answers that pass keyword matching via term overlap). The resulting C4g–C1 delta under LLM judge (+28.5 pp) is lower than the keyword delta (+47.2 pp) but remains substantial and directionally consistent. The per-condition asymmetry—C1 gains, C4g loses—is consistent with the physician finding that the keyword evaluator disproportionately penalizes C1 abstentions.

Per-category analysis. The LLM judge is particularly stricter on change (C4g: 62% mean score vs. 100% keyword) because it penalizes partial medication lists that contain correct change keywords but miss specific drugs. Conversely, the judge is more generous on historical (C4g: 62% vs. 60% keyword) and negation (C4g: 79% vs. 81% keyword), where its semantic understanding better captures correct answers that use varied phrasing.

Limitations. Using the same model family (Opus) for both answering and judging introduces potential self-preference bias. The judge also shows moderate agreement with the keyword evaluator ($\kappa = 0.572$), indicating they capture partially overlapping but distinct aspects of answer quality.

W Assertion Classifier Evaluation

The rule-based assertion classifier uses 122 calibrated trigger patterns extending the NegEx [24] and ConText [25] frameworks to a 7-class taxonomy (Table 19). Table 22 summarizes the pattern inventory and confidence ranges by category.

Table 22: Assertion trigger pattern inventory. Confidence ranges reflect per-trigger calibration.

Category	Patterns	Confidence	Example triggers
Absent	31	0.85–0.98	“no evidence of,” “denies,” “ruled out”
Uncertain	32	0.35–0.75	“possible,” “likely,” “consistent with”
Present	27	0.85–0.98	“confirmed,” “diagnosed with,” “positive for”
Hypothetical	11	0.25–0.35	“risk of,” “screening for,” “prophylaxis”
Family history	8	0.85–0.95	“family history of,” “maternal history”
Conditional	7	0.20–0.40	“if,” “contingent on,” “depending on”
Historical	6	0.75–0.90	“history of,” “former,” “remote history”
Total	122	0.20–0.98	

Corpus statistics. Across the 43 ClinicalBench patients, the knowledge graph contains 3,943 edges linked to clinical facts. Of these, 618 (15.7%) carry non-present assertions: 442 absent (11.2%), 94 possible (2.4%), 71 historical (1.8%), 8 conditional (0.2%), and 3 hypothetical (0.1%). At the mention level, 1,428 of 12,379 mentions (11.5%) are non-present. This non-present fraction $f_{np} = 0.157$ provides the empirical bound from Corollary 4: an assertion-blind pipeline cannot exceed $1 - f_{np} = 84.3\%$ assertion-faithful accuracy on concepts where negated or uncertain mentions are present.

Literature-grounded accuracy. No annotated assertion corpus was available for intrinsic evaluation; we instead bound expected accuracy using published benchmarks on similar rule-based classifiers. NegEx achieves 94.5% precision / 77.8% recall for negation [24]; ConText reports 93.8% precision for negation scope [25]; NegBio reaches 96.3% precision / 85.7% recall [28]. Our classifier extends these with calibrated confidence, pseudo-negation handling, and scope-aware bidirectional matching. Expected accuracy is ~ 85 – 95% for negation (the most studied and our largest non-present class at 71.5% of non-present edges) and ~ 75 – 85% for uncertainty and historical (less studied, fewer trigger patterns).

Functional evaluation. Rather than relying solely on intrinsic metrics, the C4 ablation provides a functional evaluation of assertion quality. C4 adds assertion metadata to C3’s graph without intent routing: the classifier’s output is directly consumed by the retrieval pipeline. On assertion-sensitive

categories (negation, conditional, uncertainty, family history), C4 outperforms C3 by an average of +16.5 pp, confirming that the classifier produces usable assertions for these categories. On temporal categories (historical, sequence, change, current state, duration), C4 underperforms C3 by an average of −24 pp—not because assertions are incorrect, but because uniform BFS traversal cannot exploit them effectively. The C4→C4g routing correction recovers these losses and amplifies the gains, achieving +22.2 pp overall ($p < 10^{-10}$).

X Cohort Demographics and Subgroup Analysis

The 43 ClinicalBench patients are drawn from MIMIC-IV [33], a single academic medical center (Beth Israel Deaconess Medical Center, Boston). Table 2 in the main text summarizes sex, age, race/ethnicity, and insurance distributions. The cohort skews female (60.5%), White (76.7%), and Medicare-insured (44.2%), reflecting the source institution’s patient mix.

Subgroup accuracy. Table 23 reports C1 and C4g accuracy by demographic group. Because questions are distributed unevenly across patients and group sizes are small ($n \leq 26$ patients per stratum), these comparisons are severely underpowered; we report them for transparency, not for drawing subgroup conclusions. No statistically significant interaction between demographic group and condition was observed, but the analysis cannot rule out meaningful effect modification.

Table 23: ClinicalBench accuracy by demographic subgroup (underpowered; reported for transparency). n_q = number of questions in each stratum.

Characteristic	Group	n_{pts}	n_q	C1 (%)	C4g (%)
Sex	Female	26	249	22.5	69.9
	Male	17	151	20.5	67.5
Age	<65	28	258	20.5	69.4
	≥65	15	142	23.9	68.3
Race	White	33	290	22.4	71.7
	Non-White	10	110	20.0	61.8

Generalizability limitations. The single-site, predominantly White cohort limits external validity. MIMIC-IV emergency department notes are heavily templated; performance on narrative-heavy specialties (psychiatry, palliative care) or community hospital documentation styles is unknown. Multi-site validation with demographically diverse cohorts is needed before deployment claims can be made.

Y Ethics Statement

This work uses de-identified clinical data from MIMIC-IV [33] under a PhysioNet Credentialed Health Data Use Agreement. No patient re-identification was attempted. The system is designed for clinical decision *support*, not autonomous clinical decision-making.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: **Yes.**

Justification: The abstract and introduction state four specific contributions (Section 1), each supported by empirical evidence in Sections 5.1–5.4. SliceBench claims include bootstrap CIs; ClinicalBench claims use within-run paired comparisons (Section 4.3). Claims are scoped with “to our knowledge” qualifiers where appropriate.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

806 Answer: **Yes.**

807 Justification: Section 6 explicitly discusses limitations including small sample sizes (6 pa-
808 tients for SliceBench, 2 per tier), single-site data (MIMIC-IV), the B2→B3 confidence inter-
809 val crossing zero, model-dependent effects, and physician pilot adjudication ($n = 30$, Sec-
810 tion 5.5).

811 **3. Theory assumptions and proofs**

812 Question: For each theoretical result, does the paper provide the full set of assumptions and a
813 complete (correct) proof?

814 Answer: **Yes.**

815 Justification: Proposition 3 (assertion entropy loss) and Corollary 4 (faithfulness bound) in
816 Appendix B state all assumptions explicitly. These are information-theoretic characterizations
817 with all steps included.

818 **4. Experimental result reproducibility**

819 Question: Does the paper fully disclose all the information needed to reproduce the main
820 experimental results of the paper to the extent that it affects the main claims and/or conclusions
821 of the paper?

822 Answer: **Yes.**

823 Justification: Appendix P reports exact model versions (MedGemma 27B, Claude Sonnet 4.5,
824 Claude Opus 4.6), temperature settings (0), bootstrap parameters ($n = 2,000$, seed 42, BCa
825 method), and data source (MIMIC-IV v1.5.0 under PhysioNet DUA). Section 4 describes
826 question generation and evaluation protocols.

827 **5. Open access to data and code**

828 Question: Does the paper provide open access to the data and code, with sufficient instructions
829 to faithfully reproduce the main experimental results?

830 Answer: **Yes.**

831 Justification: A standalone reproducibility package is included as supplementary material
832 (epikg-benchmark/) containing all 400 questions, raw predictions for all condition×model
833 combinations (7 Opus conditions, 2 MedGemma, 2 GPT-OSS), and the deterministic keyword
834 evaluator. All main-text accuracy numbers are independently reproducible. MIMIC-IV source
835 notes require separate PhysioNet credentialed access. Full system code will be released upon
836 acceptance.

837 **6. Experimental setting/details**

838 Question: Does the paper specify all the training and test details necessary to understand the
839 results?

840 Answer: **Yes.**

841 Justification: Section 4 specifies benchmark design, question counts, condition definitions,
842 patient selection criteria, and evaluation methodology. Appendix P provides model versions
843 and computational requirements.

844 **7. Experiment statistical significance**

845 Question: Does the paper report error bars suitably and correctly?

846 Answer: **Yes.**

847 Justification: Table 6 reports BCa bootstrap 95% CIs for all conditions. Per-tier results (Fig-
848 ure 3) are clearly labeled as point estimates without per-tier CIs due to small tier-level sample
849 sizes ($n = 2$ patients per tier). The overall B2→B3 CI crossing zero is explicitly noted.

850 8. Experiments compute resources

851 Question: For each experiment, does the paper provide sufficient information on the computer
852 resources needed to reproduce the experiments?

853 Answer: **Yes.**

854 Justification: Appendix P reports computational requirements: ClinicalBench Opus ≈ 22 min
855 per condition (API), MedGemma ≈ 3.6 hours (single GPU), SliceBench ≈ 2 hours (API).

856 9. Code of ethics

857 Question: Does the research conducted in the paper conform, in every respect, with the
858 NeurIPS Code of Ethics?

859 Answer: **Yes.**

860 Justification: The study uses de-identified clinical data under an approved data use agreement
861 (Appendix Y). No patient re-identification was attempted. The system is designed for clinical
862 decision *support*, not autonomous decision-making.

863 10. Broader impacts

864 Question: Does the paper discuss both potential positive societal impacts and negative societal
865 impacts of the work?

866 Answer: **Yes.**

867 Justification: Section 6 discusses the potential for improved clinical decision support (positive)
868 and risks of over-reliance on automated systems, context dilution, and the need for physician
869 oversight (negative). Appendix Y addresses data ethics.

870 11. Safeguards

871 Question: Does the paper describe safeguards that have been put in place for responsible
872 release of data or models?

873 Answer: **Yes.**

874 Justification: The system operates on de-identified data under PhysioNet DUA. It is designed
875 as a decision support tool requiring physician oversight. Code release will include usage
876 guidelines emphasizing the system is not validated for autonomous clinical use.

877 12. Licenses for existing assets

878 Question: Are the creators or original owners of assets used in the paper properly credited and
879 are the license and terms of use explicitly mentioned and properly respected?

880 Answer: **Yes.**

881 Justification: MIMIC-IV is cited [33] and used under PhysioNet Credentialed Health Data
882 Use Agreement. All benchmark frameworks and models are cited. OMOP CDM is an open
883 standard.

884 13. New assets

885 Question: Are new assets introduced in the paper well documented and is the documentation
886 provided alongside the assets?

887 Answer: **Yes.**

888 Justification: ClinicalBench and SliceBench are new benchmarks fully documented in Sec-
889 tion 4, with question generation procedures, category definitions (Appendix Q), and evalua-
890 tion protocols. Templates and scoring scripts will accompany the code release.

891 **14. Crowdsourcing and research with human subjects**

892 Question: For crowdsourcing experiments and research with human subjects, does the paper
893 include the full text of instructions given to participants and screenshots?

894 Answer: **N/A.**

895 Justification: No crowdsourcing was used. Physician adjudication (Appendix N) involves
896 expert review by co-authors, not human subjects research.

897 **15. IRB approvals or equivalent**

898 Question: Does the paper describe potential risks and does it include the IRB approval number
899 if applicable?

900 Answer: **N/A.**

901 Justification: MIMIC-IV is a de-identified public dataset; use under PhysioNet DUA does not
902 require separate IRB approval per PhysioNet policy.

903 **16. Declaration of LLM usage**

904 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-
905 standard component of the core methods?

906 Answer: **Yes.**

907 Justification: LLMs are core experimental components, not authoring aids. Claude Opus 4.6
908 is the primary answering model for ClinicalBench (Section 5.1); Claude Sonnet 4.5 answers
909 SliceBench questions while Opus 4.6 judges (Section 5.4); MedGemma 27B provides cross-
910 model validation (Section 5.3). Model versions, temperatures, and inference details are in
911 Appendix P.